

Starting April 29, 2025, Gemini 1.5 Pro and Gemini 1.5 Flash models are not available in projects that have no prior usage of these models, including new projects. For details, see [Model versions and lifecycle](#) (/vertex-ai/generative-ai/docs/learn/model-versions#legacy-stable).

Lyria API

[Release Notes](#)

Lyria is a new foundation model for high-quality audio generation, capable of creating diverse soundscapes and musical pieces from text prompts. Lyria enables users to generate high-quality instrumental music from text prompts.

To explore this model in the console, see the Lyria model card in the Model Garden (accessible using the Media Studio tab).

[Try Lyria on Vertex AI \(Vertex AI Studio\)](#) (<https://console.cloud.google.com/vertex-ai/studio/media/generate;tab=music>)

[Try Lyria in a Colab](#) (https://github.com/GoogleCloudPlatform/generative-ai/blob/main/audio/music/getting-started/lyria2_music_generation.ipynb)

Supported Models

The Lyria API supports the following model:

- lyria-002

HTTP request

```
curl -X POST \
-H "Authorization: Bearer $(gcloud auth print-access-token)" \
-H "Content-Type: application/json" \
https://LOCATION[-aiplatform.googleapis.com/v1/projects/](https://-aiplatform.googleapis.com/v1/projects/)
-d '{
  "instances": [
    {
      "prompt": "string",
      "negative_prompt": "string", // Optional
      "seed": 0 // Optional. Cannot be used with sample_count.
    }
  ],
  "parameters": {
    "sample_count": 1 // Optional. Cannot be used with seed.
  }
}'
```

Use the following parameters for the Lyria model. For more information, see the Lyria Model Garden card details.

Parameter	
prompt	string

Parameter	
(in <code>instances</code> object)	Required. The text description in US English (en-us) of the audio to generate. Example: <i>"An energetic electronic dance track with a fast tempo."</i>
<code>negative_prompt</code>	<code>string</code>
(in <code>instances</code> object)	Optional. A description of what to exclude from the generated audio. Example: <i>"vocals, slow tempo"</i>
<code>seed</code>	<code>integer</code>
(in <code>instances</code> object)	Optional. A seed for deterministic generation. If provided, the model will attempt to produce the same audio given the same prompt and other parameters. Cannot be used with <code>sample_count</code> in the same request. Example: 12345
<code>sample_count</code>	<code>integer</code>
(in <code>parameters</code> object)	Optional. The number of audio samples to generate. Default is 1 if not specified and seed is not used. Cannot be used with <code>seed</code> in the same request. Example: 2

Sample request

Use the following request to generate instrumental music from a text prompt:

Text-to-music generation request

curlJSON (#json)JSON with sample_count...
(#curl)

```
curl -X POST \
-H "Authorization: Bearer $(gcloud auth print-access-token)" \
-H "Content-Type: application/json" \
https://us-central1-aiplatform.googleapis.com/v1/projects/PROJECT_ID/locations/us-central1/publishers,
-d '{
  "instances": [
    {
      "prompt": "A calm acoustic folk song with a gentle guitar melody and soft strings.",
      "negative_prompt": "drums, electric guitar",
      "seed": 98765
    }
  ],
  "parameters": {}
}
```

Response body

A successful request returns a JSON object containing the generated audio data. Each generated audio clip is 30 seconds long and provided as a WAV audio file at a 48kHz sample rate.

```
{
  "predictions": [
    {
      "audioContent": "BASE64_ENCODED_WAV_STRING_SAMPLE_1",
      "mimeType": "audio/wav"
    }
    // Additional audio samples will be listed here if sample_count > 1
    // e.g.,
    // {"audioContent": "BASE64_ENCODED_WAV_STRING_SAMPLE_2",
    //   "mimeType": "audio/wav"
    // }
  ],
  "deployedModelId": "xxxxxxxxxxxxxxxx", // Actual ID may vary based on deployment
  "model": "projects/PROJECT_ID /locations/LOCATION /publishers/google/models/lyria-002",
  "modelDisplayName": "Lyria 2"
}
```

Response element	
predictions	array An array of generated audio samples. Each object in the array represents one audio clip.
predictions[].audioContent	string Base64-encoded string of the generated WAV audio data.
predictions[].mimeType	string The MIME type of the audio data. For Lyria, this is "audio/wav".
deployedModelId	string The ID of the deployed model that processed the request (if applicable for the endpoint type).
model	string The full resource name of the model that processed the request.
modelDisplayName	string The display name of the model.

Best practices and limitations

Refer to the Lyria Model Card for detailed best practices on prompting, language support (US English only for prompts), generation times, output format (WAV, 48 kHz, 30s instrumental clips), safety measures, and deployment information.

Key points:

- **Detailed Prompts:** Generally lead to better audio.
- **Specify:** Genre, mood, instrumentation, tempo.
- **Negative Prompting:** Use negative_prompt to exclude elements.

- **Output:** 30-second WAV audio clips, 48 kHz, instrumental only.
- **Safety:** Content safety filters, recitation checking, artist intent checks, and SynthID watermarking are applied.

Pricing

Lyria 2 usage is priced at \$0.06 per 30 seconds of output music generated. For more details, see [Vertex AI Pricing](https://cloud.google.com/vertex-ai/pricing) (<https://cloud.google.com/vertex-ai/pricing>).

More information

- Learn more about [Generative AI on Vertex AI](https://cloud.google.com/vertex-ai/docs/generative-ai/) (<https://cloud.google.com/vertex-ai/docs/generative-ai/>).
- For an overview of Lyria, refer to its model card available in the Model Garden (Media Studio).

What's next

- Try out Lyria in the [Vertex AI Studio](https://console.cloud.google.com/vertex-ai/studio/media/generate;tab=music) (<https://console.cloud.google.com/vertex-ai/studio/media/generate;tab=music>).
- Review the [Google Cloud Service Terms](https://cloud.google.com/terms) (<https://cloud.google.com/terms>).
- Read the [Additional Terms for Generative AI Products](https://cloud.google.com/terms/generative-ai) (<https://cloud.google.com/terms/generative-ai>).

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 4.0 License](https://creativecommons.org/licenses/by/4.0/) (<https://creativecommons.org/licenses/by/4.0/>), and code samples are licensed under the [Apache 2.0 License](https://www.apache.org/licenses/LICENSE-2.0) (<https://www.apache.org/licenses/LICENSE-2.0>). For details, see the [Google Developers Site Policies](https://developers.google.com/site-policies) (<https://developers.google.com/site-policies>). Java is a registered trademark of Oracle and/or its affiliates.

Last updated 2025-06-06 UTC.